

# Data to a network visualisation

*Mike Spencer*

*14 March 2018*

## Intro

This document has been written in R and accompanies the live coding part of the network analysis workshop. In this document you'll see a mixture of code and output. Hopefully it'll be easy to tell these apart! To help, lines of output begin with `##`.

## Packages

- Install only once
- Load into session with `library()`

```
# install.packages("tidyverse")
# install.packages("igraph")
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.2.1 --
```

```
## ✓ ggplot2 2.2.1      ✓ purrr  0.2.4
## ✓ tibble  1.4.2      ✓ dplyr  0.7.4
## ✓ tidyr   0.8.0      ✓ stringr 1.2.0
## ✓ readr   1.1.1      ✓ forcats 0.2.0
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(igraph)
```

```
##
```

```
## Attaching package: 'igraph'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
##   as_data_frame, groups, union
```

```
## The following objects are masked from 'package:purrr':
```

```
##
```

```
##   compose, simplify
```

```
## The following object is masked from 'package:tidyr':
```

```
##
```

```
##   crossing
```

```
## The following object is masked from 'package:tibble':
```

```
##
```

```
##   as_data_frame
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##   decompose, spectrum
```

```
## The following object is masked from 'package:base':
##
##      union
```

## Reading data

We can read data from local files, but as you'll have seen in the workshop we can also read files from a web address.

```
# Reads and outputs to console
read_csv("../data/SNA_anon_delegates.csv")
```

```
## Parsed with column specification:
## cols(
##   Timestamp = col_character(),
##   `Your name` = col_character(),
##   `Your affiliation` = col_character(),
##   `Your seniority` = col_character(),
##   `Primary software expertise` = col_character(),
##   `I really want to learn...` = col_character()
## )

## # A tibble: 27 x 6
##   Timestamp      `Your name` `Your affiliation` `Your seniority`
##   <chr>          <chr>      <chr>             <chr>
## 1 02/03/2018 21:44:54 Person 1    Research - AVS    Early career
## 2 02/03/2018 21:47:45 Person 2    Research - LEES   Early career
## 3 02/03/2018 23:03:18 Person 3    Research - FFS    Student
## 4 03/03/2018 04:53:02 Person 4    Research - FFS    Early career
## 5 03/03/2018 10:26:59 Person 5    Research - AVS    Student
## 6 03/03/2018 10:32:22 Person 6    Research - FFS    Early career
## 7 03/03/2018 11:56:24 Person 7    Research - AVS    Student
## 8 03/03/2018 13:30:19 Person 8    Research - CSS    Early career
## 9 03/03/2018 14:39:47 Person 9    Research - LEES   Mid career
## 10 04/03/2018 08:48:31 Person 10   Research - FFS    Senior
## # ... with 17 more rows, and 2 more variables: `Primary software
## #   expertise` <chr>, `I really want to learn...` <chr>
```

```
# Reads and assigns to object df
df = read_csv("../data/SNA_anon_delegates.csv")
```

```
## Parsed with column specification:
## cols(
##   Timestamp = col_character(),
##   `Your name` = col_character(),
##   `Your affiliation` = col_character(),
##   `Your seniority` = col_character(),
##   `Primary software expertise` = col_character(),
##   `I really want to learn...` = col_character()
## )
```

## Cleaning data

This section prepares the dataset a little for exploratory analysis. It's worth noting here, that I've avoided tidying the `want_to_learn` columns at this point.

```
# Shortening and removing spaces from column names
colnames(df) = c("timestamp", "name", "affiliation", "seniority", "expertise", "want_to_learn")

# Reducing the long other answers
df$expertise[df$expertise=="Network analysis software like biolayout/Miru"] = "Biolayout"
df$expertise[df$expertise=="Excel (advanced)"] = "MS/Libre/Open office"
```

## Selecting columns

Particularly if we're working with large datasets, it can be useful to pull out the columns we're interested in.

```
# Data followed by columns we want
select(df, name, affiliation, seniority, expertise)

## # A tibble: 27 x 4
##   name      affiliation      seniority  expertise
##   <chr>      <chr>          <chr>      <chr>
## 1 Person 1 Research - AVS Early career MaxQDA
## 2 Person 2 Research - LEES Early career FORTRAN
## 3 Person 3 Research - FFS Student      R
## 4 Person 4 Research - FFS Early career SQL
## 5 Person 5 Research - AVS Student      R
## 6 Person 6 Research - FFS Early career R
## 7 Person 7 Research - AVS Student      MS/Libre/Open office
## 8 Person 8 Research - CSS Early career R
## 9 Person 9 Research - LEES Mid career  Stata
## 10 Person 10 Research - FFS Senior      Biolayout
## # ... with 17 more rows

# Or data followed by columns we don't want
select(df, -timestamp, -want_to_learn)

## # A tibble: 27 x 4
##   name      affiliation      seniority  expertise
##   <chr>      <chr>          <chr>      <chr>
## 1 Person 1 Research - AVS Early career MaxQDA
## 2 Person 2 Research - LEES Early career FORTRAN
## 3 Person 3 Research - FFS Student      R
## 4 Person 4 Research - FFS Early career SQL
## 5 Person 5 Research - AVS Student      R
## 6 Person 6 Research - FFS Early career R
## 7 Person 7 Research - AVS Student      MS/Libre/Open office
## 8 Person 8 Research - CSS Early career R
## 9 Person 9 Research - LEES Mid career  Stata
## 10 Person 10 Research - FFS Senior      Biolayout
## # ... with 17 more rows
```

## Filter by row value

What if we're not interested in every observation? Maybe we only want to look at those respondents from Land Economy, or find those with expertise in R.

```
# Single filter
```

```
filter(df, affiliation=="Research - LEES")
```

```
## # A tibble: 9 x 6
```

##	timestamp	name	affiliation	seniority	expertise	want_to_learn
##	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
## 1	02/03/2018 21:47:45	Perso~	Research --	Early ca~	FORTTRAN	R
## 2	03/03/2018 14:39:47	Perso~	Research --	Mid care~	Stata	R, Stata
## 3	04/03/2018 13:03:34	Perso~	Research --	Student	R	R, Python
## 4	04/03/2018 21:56:33	Perso~	Research --	Mid care~	MS/Libre~	R
## 5	05/03/2018 09:25:32	Perso~	Research --	Student	MS/Libre~	R, QGIS
## 6	05/03/2018 09:55:59	Perso~	Research --	Early ca~	R	SQL, R, Pyth~
## 7	05/03/2018 17:10:43	Perso~	Research --	Student	MS/Libre~	SQL, R, SPSS~
## 8	06/03/2018 17:10:11	Perso~	Research --	Mid care~	R	R, Python, A~
## 9	08/03/2018 01:25:19	Perso~	Research --	Student	None	R

```
# Exclude
```

```
filter(df, expertise!="R")
```

```
## # A tibble: 16 x 6
```

##	timestamp	name	affiliation	seniority	expertise	want_to_learn
##	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
## 1	02/03/2018 21:44:54	Pers~	Research --	Early ca~	MaxQDA	R, SPSS
## 2	02/03/2018 21:47:45	Pers~	Research --	Early ca~	FORTTRAN	R
## 3	03/03/2018 04:53:02	Pers~	Research --	Early ca~	SQL	R
## 4	03/03/2018 11:56:24	Pers~	Research --	Student	MS/Libre~	R
## 5	03/03/2018 14:39:47	Pers~	Research --	Mid care~	Stata	R, Stata
## 6	04/03/2018 08:48:31	Pers~	Research --	Senior	Biolayout	R
## 7	04/03/2018 21:56:33	Pers~	Research --	Mid care~	MS/Libre~	R
## 8	05/03/2018 09:25:32	Pers~	Research --	Student	MS/Libre~	R, QGIS
## 9	05/03/2018 11:46:48	Pers~	Corporate/~	Mid care~	CBS	VBA
## 10	05/03/2018 17:10:43	Pers~	Research --	Student	MS/Libre~	SQL, R, SPSS~
## 11	05/03/2018 17:42:18	Pers~	Consulting	Senior	CBS	R
## 12	06/03/2018 09:09:50	Pers~	Consulting	Senior	MS/Libre~	SQL
## 13	07/03/2018 16:44:49	Pers~	Consulting	Mid care~	MS/Libre~	R, Excel (ad~
## 14	08/03/2018 01:25:19	Pers~	Research --	Student	None	R
## 15	08/03/2018 13:15:20	Pers~	Education	Mid care~	ArcGIS	R
## 16	15/03/2018 10:56:00	Pers~	Research --	Student	ArcGIS	R

```
# Multiple filters? use & (and) or | (or)
```

```
filter(df, affiliation=="Research - LEES" & expertise=="R")
```

```
## # A tibble: 3 x 6
```

##	timestamp	name	affiliation	seniority	expertise	want_to_learn
##	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
## 1	04/03/2018 13:03:34	Perso~	Research --	Student	R	R, Python
## 2	05/03/2018 09:55:59	Perso~	Research --	Early ca~	R	SQL, R, Pyth~
## 3	06/03/2018 17:10:11	Perso~	Research --	Mid care~	R	R, Python, A~

```
# With a pipe
```

```
df %>%
```

```
  select(-timestamp, -want_to_learn) %>%
```

```
filter(affiliation=="Research - LEES" & expertise=="R")
```

```
## # A tibble: 3 x 4
##   name      affiliation      seniority expertise
##   <chr>      <chr>          <chr>      <chr>
## 1 Person 11 Research - LEES Student      R
## 2 Person 15 Research - LEES Early career R
## 3 Person 23 Research - LEES Mid career   R
```

```
# With numbers
# filter(df, col_num==10)
# filter(df, col_num>10)
# etc.
# Note these commented lines of filter() are not run.
```

## Summaries

We often want to summarise our data. This may be simple counts of categories, or it may be numerical methods like taking a mean. The `count` command simply counts how many of each thing occur in a column.

If we want to do more than this we can use `summarise`, but in order to do this we need to tell R how to group our data. `group_by` tells R which column(s) to group our data on. If we had already cleaned our `want_to_learn` column into a tidy format (Wickham 2014 <http://vita.had.co.nz/papers/tidy-data.pdf>), most of our examples would have needed to use `group_by`.

```
# Basic how many?
count(df, expertise)
```

```
## # A tibble: 10 x 2
##   expertise      n
##   <chr>      <int>
## 1 ArcGIS          2
## 2 Biolayout       1
## 3 CBS             2
## 4 FORTRAN         1
## 5 MaxQDA          1
## 6 MS/Libre/Open office 6
## 7 None           1
## 8 R             11
## 9 SQL            1
## 10 Stata          1
```

```
# Ordered
df %>%
  count(expertise) %>%
  arrange(n)
```

```
## # A tibble: 10 x 2
##   expertise      n
##   <chr>      <int>
## 1 Biolayout       1
## 2 FORTRAN         1
## 3 MaxQDA          1
## 4 None           1
## 5 SQL            1
```

```
## 6 Stata 1
## 7 ArcGIS 2
## 8 CBS 2
## 9 MS/Libre/Open office 6
## 10 R 11

# By more categories we can use group_by
df %>%
  group_by(seniority, expertise) %>%
  summarise(n=n()) %>%
  arrange(n)


## # A tibble: 16 x 3
## # Groups:   seniority [4]
##   seniority expertise n
##   <chr> <chr> <int>
## 1 Early career FORTRAN 1
## 2 Early career MaxQDA 1
## 3 Early career SQL 1
## 4 Mid career ArcGIS 1
## 5 Mid career CBS 1
## 6 Mid career R 1
## 7 Mid career Stata 1
## 8 Senior Biolayout 1
## 9 Senior CBS 1
## 10 Senior MS/Libre/Open office 1
## 11 Student ArcGIS 1
## 12 Student None 1
## 13 Mid career MS/Libre/Open office 2
## 14 Early career R 3
## 15 Student MS/Libre/Open office 3
## 16 Student R 7

# For a mean
# df %>%
#   group_by(seniority, expertise) %>%
#   summarise(mean_col1=mean(col1))
```

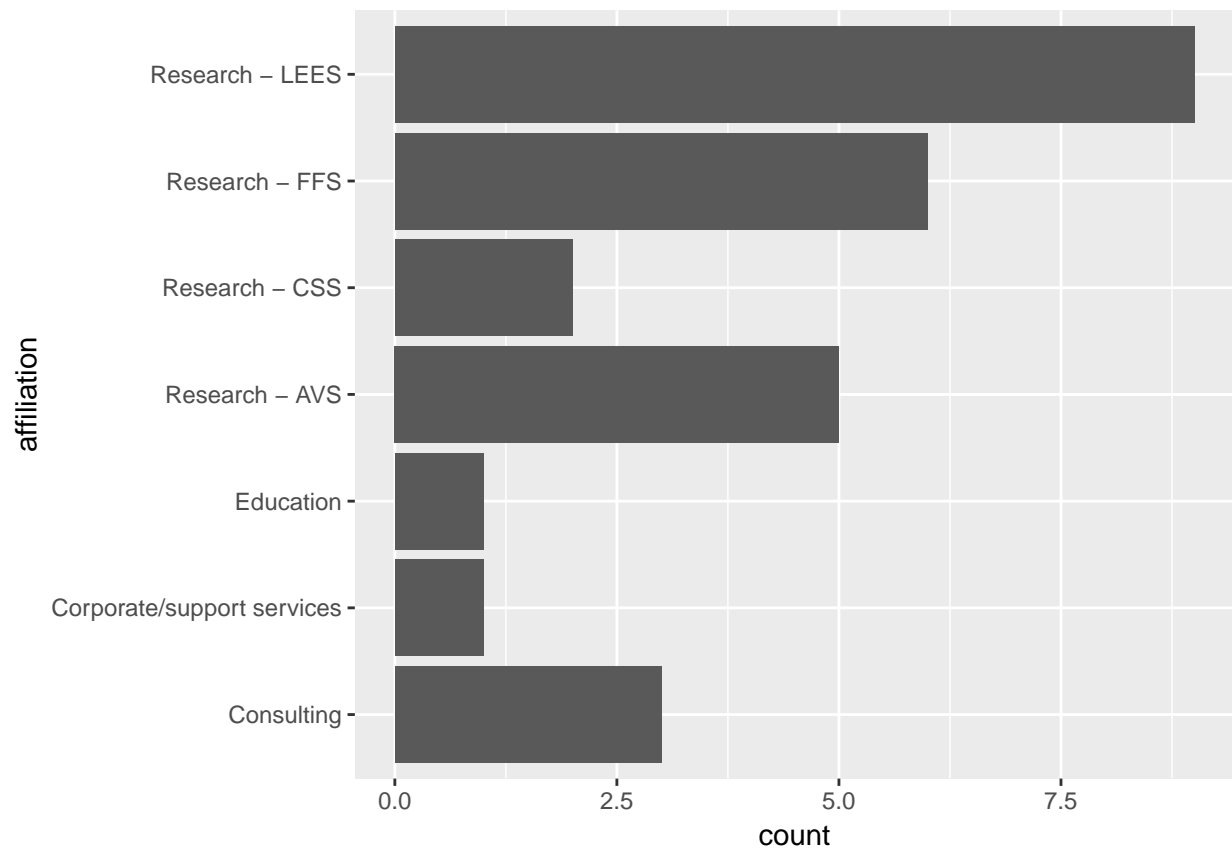
## Plots

R is *really* powerful for making plots. There are a number of ways to do this, we're going to use the `ggplot2` package. Have a look here <http://ggplot2.tidyverse.org/reference/> to give you an idea of some of the things we can do!

```
ggplot(df, aes(affiliation)) +
  geom_bar()
```

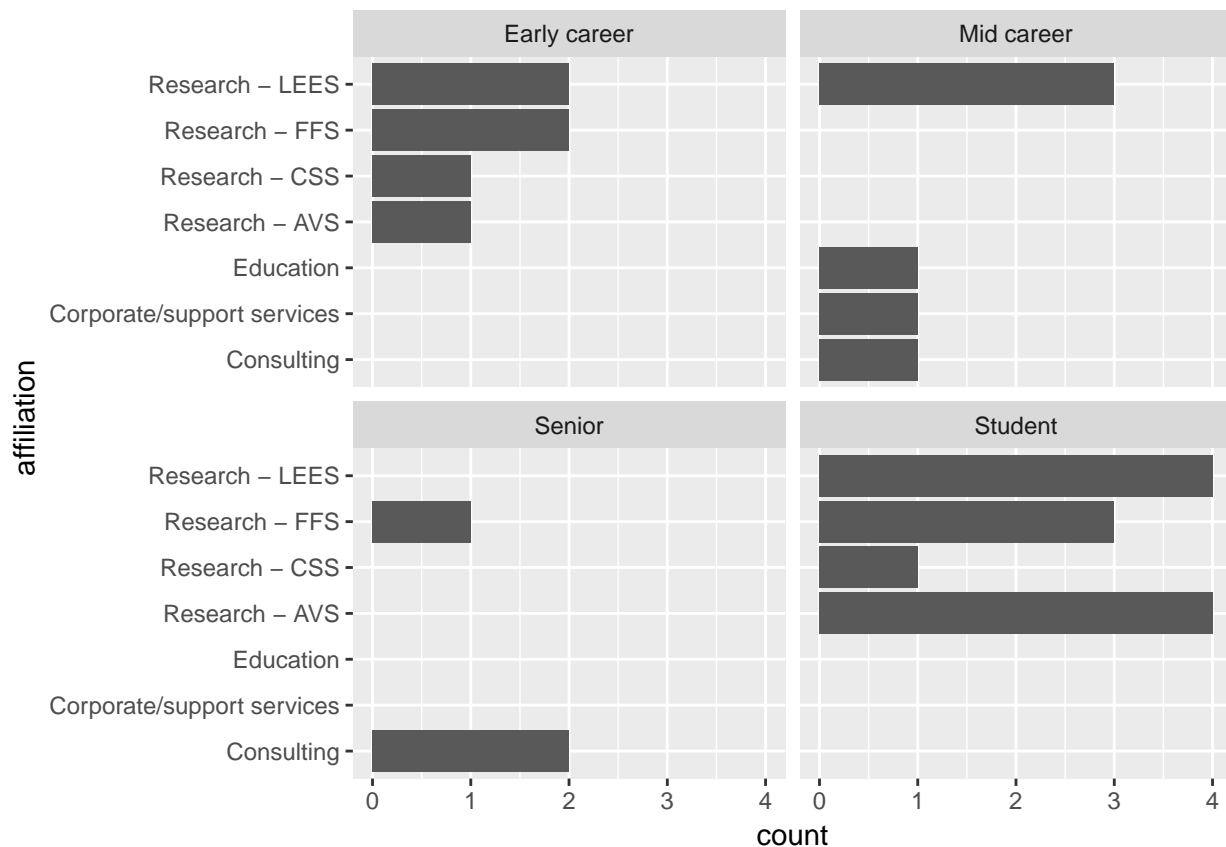


```
ggplot(df, aes(affiliation)) +  
  geom_bar() +  
  coord_flip()
```



```
ggplot(df, aes(affiliation)) +  
  geom_bar() +  
  coord_flip() +  
  facet_wrap(~ seniority)
```





## What time did you get up?

We can take the time stamps of registration and see how they spread across peoples' (self assessed) level of seniority. This is the tip of the iceberg on why data science can be considered intrusive. Note we can't really read anything into this as the sample sizes are very small.

Here we're introducing `mutate` to add extra variables.


```
# Extract the hour of registration from the timestamp column
df %>%
  select(timestamp) %>%
  mutate(hr=substr(timestamp, 12, 13))
```

```
## # A tibble: 27 x 2
##   timestamp      hr
##   <chr>         <chr>
## 1 02/03/2018 21:44:54 21
## 2 02/03/2018 21:47:45 21
## 3 02/03/2018 23:03:18 23
## 4 03/03/2018 04:53:02 04
## 5 03/03/2018 10:26:59 10
## 6 03/03/2018 10:32:22 10
## 7 03/03/2018 11:56:24 11
## 8 03/03/2018 13:30:19 13
## 9 03/03/2018 14:39:47 14
## 10 04/03/2018 08:48:31 08
## # ... with 17 more rows
```

```
# As a number
x = df %>%
  select(timestamp, seniority) %>%
  mutate(hr=as.numeric(substr(timestamp, 12, 13)))
```

```
# All registrations
ggplot(x, aes(hr)) +
  geom_histogram() +
  labs(title="Hour of day participants registered",
       x="Hour of the day",
       y="Number of participants")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

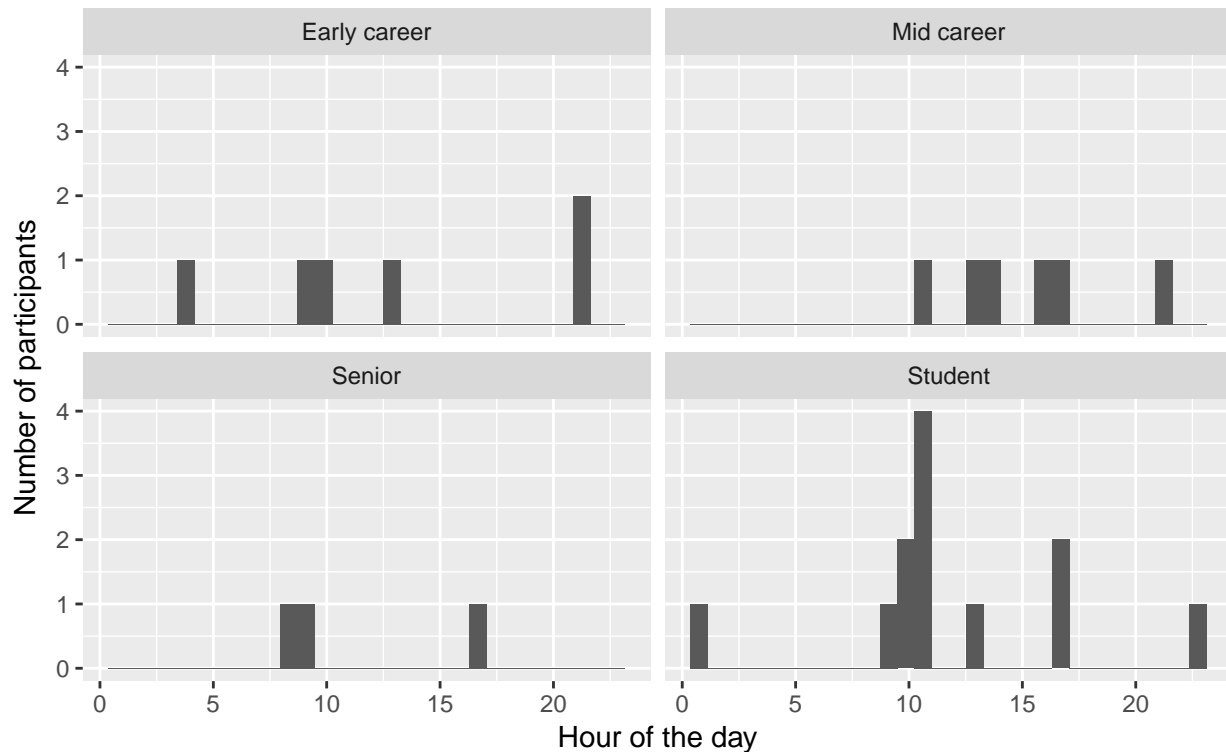


```
# Split by seniority
ggplot(x, aes(hr)) +
  geom_histogram() +
  facet_wrap(~seniority) +
  labs(title="Hour of day participants registered",
       subtitle="Split by seniority",
       x="Hour of the day",
       y="Number of participants")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

## Hour of day participants registered

Split by seniority



## Tidy data

The concept of tidy data is where each column is a variable and each row is an observation. It's worth repeating: Wickham 2014 is excellent <http://ggplot2.tidyverse.org/reference/>. Here we're going to use `str_count` to find out how many things each participant wants to learn and the split into separate columns.

```
# Maximum software types
n = str_count(df$want_to_learn, ",") %>%
  max() + 1

# Wide not tidy data
df.learning = df %>%
  select(name, want_to_learn) %>%
  separate(want_to_learn, paste0("learn_", 1:n), sep=", ", fill="right")

# Tidy data
# Double thumbs up
df.learning = df %>%
  select(name, want_to_learn) %>%
  separate(want_to_learn, paste0("learn_", 1:n), sep=", ", fill="right") %>%
  gather(ToDelete, want_to_learn, -name, na.rm=T) %>%
  select(-ToDelete)
```

## Joining data

We've now got a separate, tidy, data frame of the software each person wants to learn. As required we can join this to our original data for use.

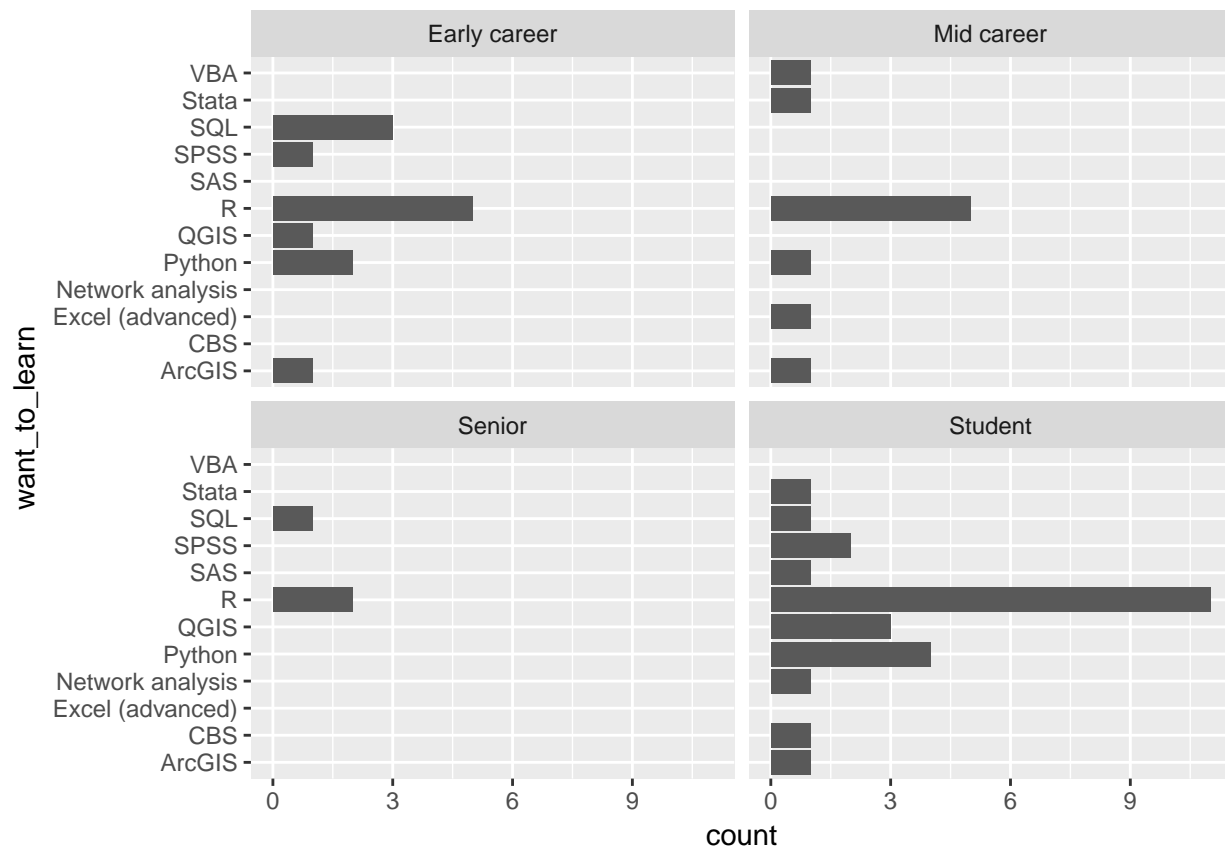
```
df %>%
  select(name, seniority) %>%
  inner_join(df.learning)

## Joining, by = "name"

## # A tibble: 52 x 3
##   name      seniority  want_to_learn
##   <chr>    <chr>    <chr>
## 1 Person 1 Early career R
## 2 Person 1 Early career SPSS
## 3 Person 2 Early career R
## 4 Person 3 Student      CBS
## 5 Person 3 Student      Stata
## 6 Person 3 Student      Python
## 7 Person 3 Student      SPSS
## 8 Person 3 Student      SAS
## 9 Person 3 Student      QGIS
## 10 Person 3 Student     ArcGIS
## # ... with 42 more rows

df %>%
  select(name, seniority) %>%
  inner_join(df.learning) %>%
  ggplot(aes(want_to_learn)) +
  geom_bar() +
  facet_wrap(~ seniority) +
  coord_flip()

## Joining, by = "name"
```



## Basic network graph

Show me the money! I know, the above doesn't look like network analysis at all, but its usefulness will hopefully become apparent.

This section moves on a lot from the earlier one. I would love there to be time to explain this code in detail, but we'll have to save that for another workshop.

The example below creates a bipartite graph, but with nodes showing the two tiers. In this case we're using people and seniority for our nodes/vertices.


```
# Make a data frame of edges
df.edges = df %>%
  select(name, seniority)

# Make a data frame of vertices
# First create a vectors of unique people and seniority levels
x = data.frame(name=unique(df$seniority), size=20)
y = data.frame(name=df$name, size=5)

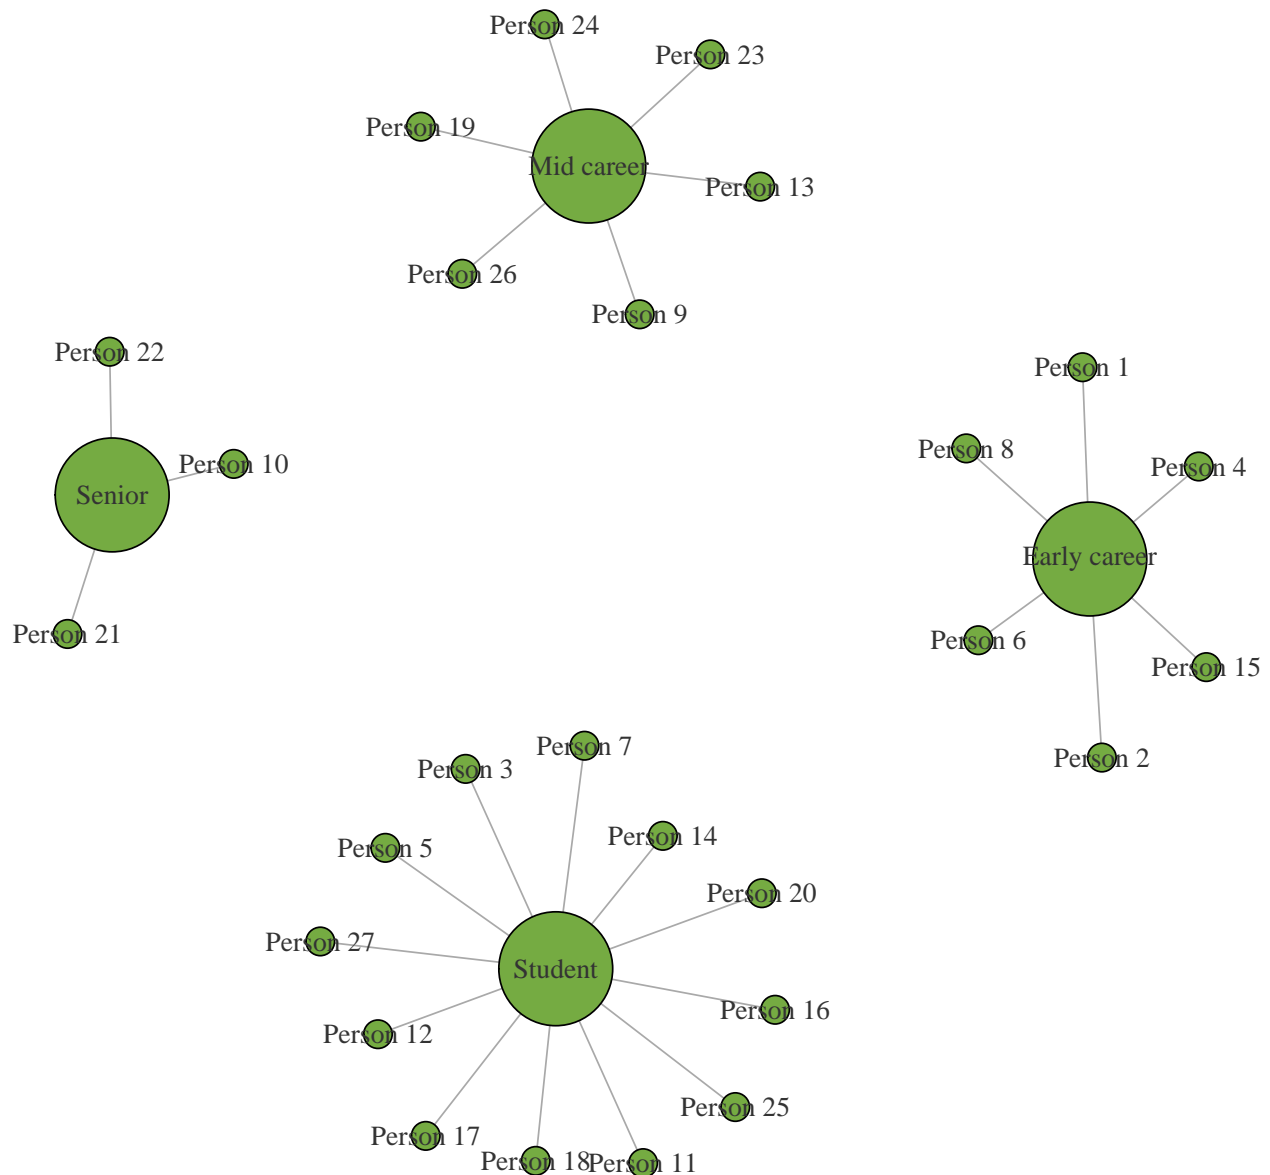
# Join these vectors together
df.vertices = rbind(x, y)

# Turn these into a graph data frame
df_graph = graph.data.frame(df.edges,
                             df.vertices,
                             directed=F)
```

```
# Plot our first graph!
plot(df_graph)
```



```
# Maybe different colours?
x = data.frame(name=unique(df$seniority), color="#75ab42", size=20)
y = data.frame(name=df$name, color="#75ab42", size=5)
df.vertices = rbind(x, y)
df_graph = graph.data.frame(df.edges,
                             df.vertices,
                             directed=F)
plot(df_graph,
     vertex.label.color="#333333")
```



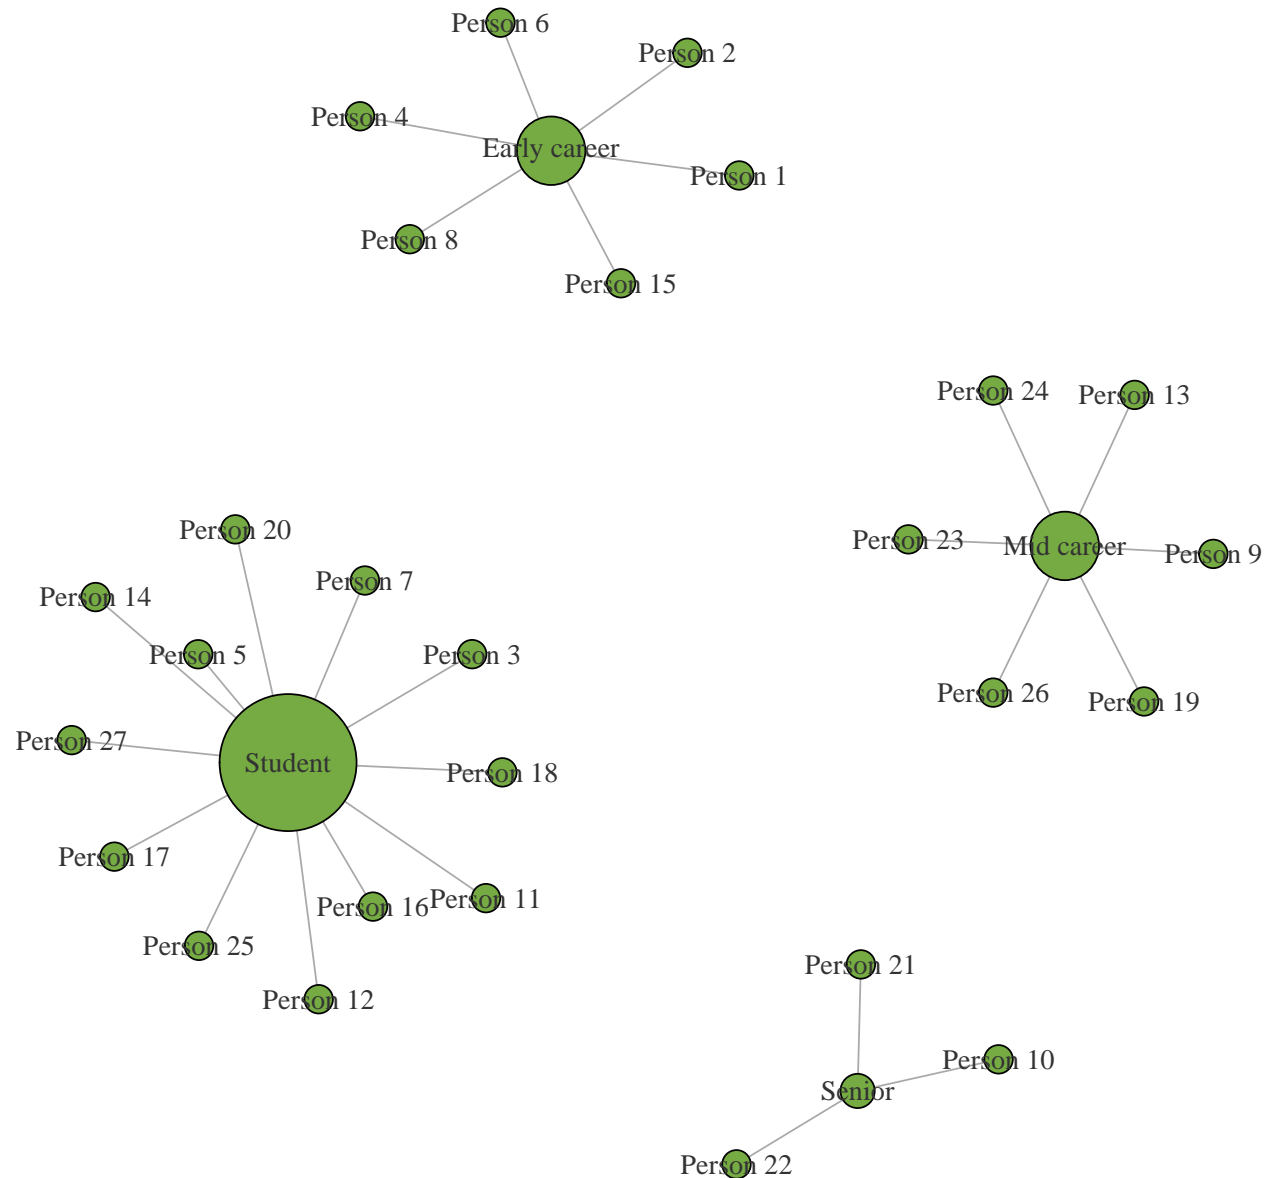
## Vertex size

But really, we might like our network to use parameters, or derived parameters to influence the way it looks. We can use the skills we learned during the earlier sections to do this.

```
# Get count of people
x = df %>%
  count(seniority) %>%
  mutate(name=seniority,
         color="#75ab42",
         size = n * 2) %>%
  select(-n, -seniority)

y = data.frame(name=df$name, color="#75ab42", size=5)
df.vertices = rbind(x, y)
df_graph = graph.data.frame(df.edges,
```

```
df.vertices,
directed=F)
plot(df_graph,
vertex.label.color="#333333")
```



## Edge weight

We can also change our edge weight to show a variable. In this example I'm using the count of different software types someone wants to learn to weight the edge.

```
# Count software types for each person
# Make into our edges
df.edges = df %>%
  select(name, seniority, want_to_learn) %>%
  mutate(width=str_count(df$want_to_learn, ", ") + 1,
```




```

        color="#004b23") %>%
select(-want_to_learn)

# Vertices
x = df %>%
  count(seniority) %>%
  mutate(name=seniority, color="#75ab42", size = n * 2) %>%
  select(-n, -seniority)
y = data.frame(name=df$name, color="#75ab42", size=5)
df.vertices = rbind(x, y)
df_graph = graph.data.frame(df.edges,
                             df.vertices,
                             directed=F)

plot(df_graph,
      vertex.label.color="#333333")

```



## Putting these ideas together

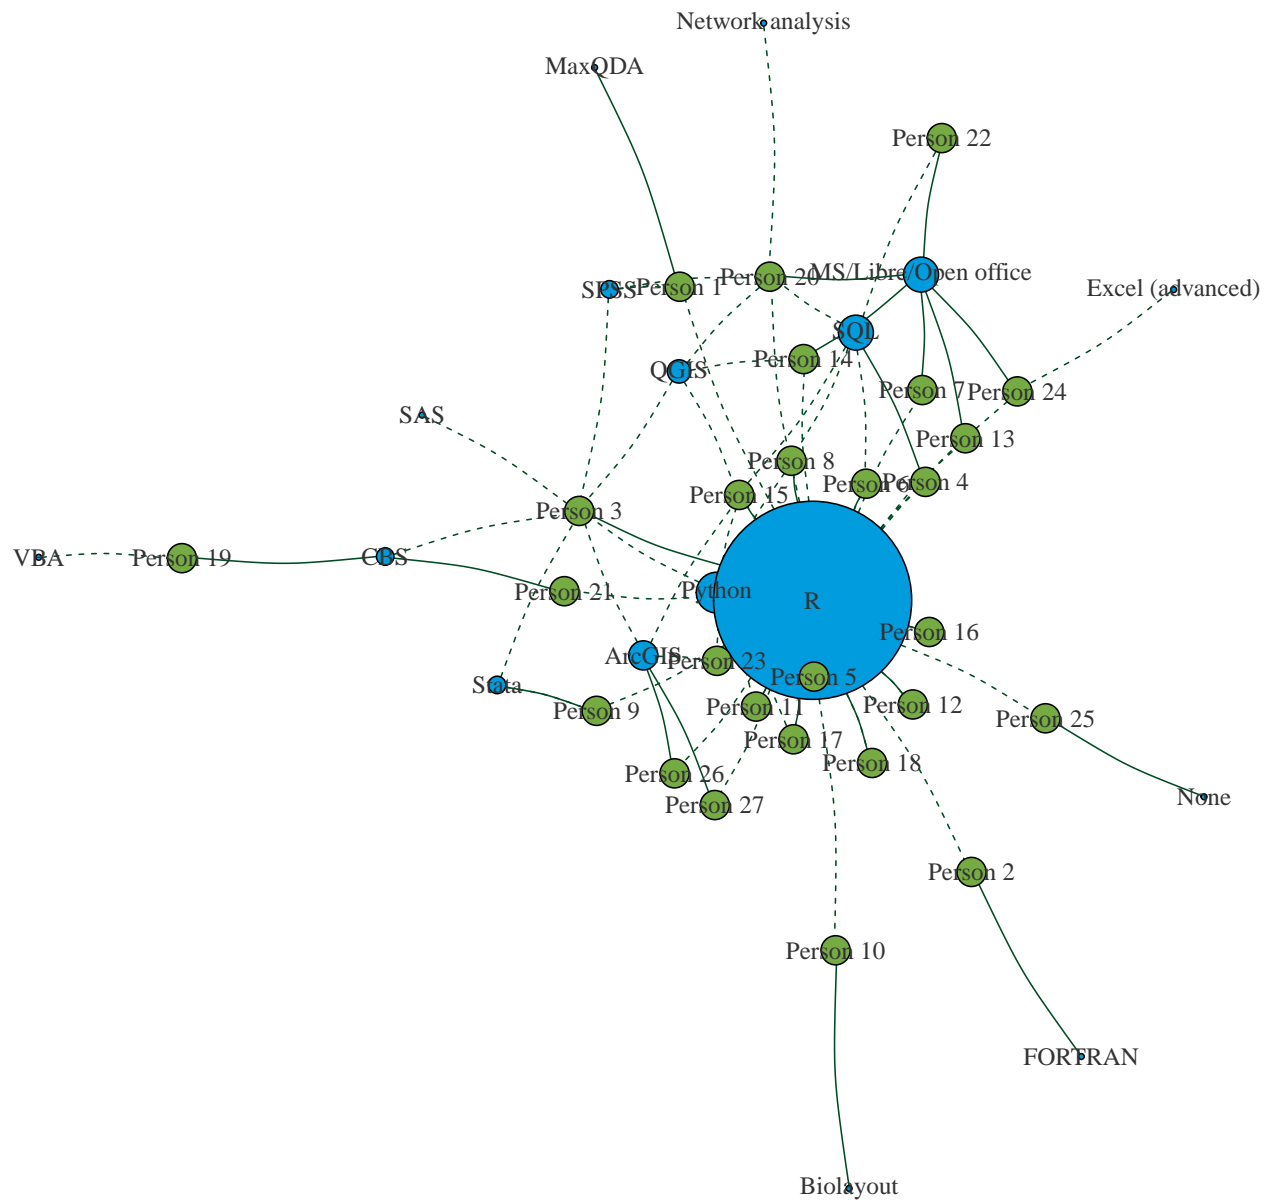
We've some potentially useful data available here. Can we use it to make a diagram of where to go for help?

```
# Want to learn edges
x = df.learning %>%
  select(name, want_to_learn) %>%
  mutate(software=want_to_learn,
         lty=2,
         width=1,
         color="#004b23") %>%
  select(-want_to_learn)
# Expertise edges
y = df %>%
  select(name, expertise) %>%
  mutate(software=expertise,
         lty=1,
         width=1,
         color="#004b23") %>%
  select(-expertise)

df.edges = rbind(x, y)

# Vertices
x = df.edges %>%
  count(software) %>%
  mutate(name=software, color=rgb(0/255, 156/255, 222/255 ), size = n) %>%
  select(-n, -software)
y = data.frame(name=df$name, color="#75ab42", size=5)
df.vertices = rbind(x, y)
df_graph = graph.data.frame(df.edges,
                           df.vertices,
                           directed=F)

par(mar=c(0.5, 0.5, 0.5, 0.5))
plot(df_graph,
     vertex.label.color="#333333",
     edge.curved=.1)
```



```
par(mar=c(5, 4, 4, 2) + 0.1)
```

Clearly more work to do to get the sizing right!